



Observations/Results

Experimented with few Openmp programs, varying the number of threads, and the scheduling pattern.



Varying the number of threads



Serial PI Program

```
static long num_steps = 100000;
double step;
int main ()
{
    int i; double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```

Tried out 3 ways to parallelize - Program 1

```
#pragma omp parallel
{
    int i;
    double sum;
    double x;
    int id = omp_get_thread_num();
    int num_threads = omp_get_num_threads();
    if(id == 0)
        nthreads = num_threads;
    for (i=id, sum = 0.0; i< num_steps; i+=num_threads){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    sum = sum * step;
#pragma omp atomic
    pi += sum;
}
```

'Atomic' section provides mutual exclusion for the threads.

Performance



Optimum :
39 threads

Program 2


```
65 #pragma omp parallel
66 {
67     int i;
68     int id = omp_get_thread_num();
69     int numthreads = omp_get_num_threads();
70     // printf(" num_threads = %d", numthreads);
71     double x;
72
73     sum[id] = 0.0;
74
75     if (id == 0) {
76         nthreads = numthreads;
77     }
78
79     for (i=id; i< num_steps; i+=numthreads){
80         x = (i+0.5)*step;
81         sum[id] = sum[id] + 4.0/(1.0+x*x);
82     }
83     //printf("Fdf");
84 }
85
86 for(full_sum = 0.0, i=0; i<nthreads; i++)
87     full_sum += sum[i];
88
89 pi = step * full_sum;
```

Here, an array is used for the variable 'sum'.

Array elements are continuous in memory, so they will share the same cache line.

In this program, there is bad cache-hit ratio, performance is not as good as before.

Performance



Optimum :
57 threads


Program-3

```
33  #pragma omp parallel
34  {
35      int i;
36      double x;
37      int id = omp_get_thread_num();
38      int num_threads = omp_get_num_threads();
39      if(id == 0)
40          nthreads = num_threads;
41      for (i=id, sum[id][0] = 0.0; i< num_steps; i+=num_threads){
42          x = (i+0.5)*step;
43          // This stores the partial sum from each thread
44          sum[id][0] = sum[id][0] + 4.0/(1.0+x*x);
45      }
46  }
47
48  int i;
49  double pi =0.0;
50  for(i = 0; i < nthreads;i++)
51      pi += sum[i][0] * step;
52
```

Here, we padded the sum array. Now it is a double array so each sum value is in a different cache line.


Performance much better than before version, which had a bad cache-hit ratio.

Performance



Optimum :
38 threads

Comparison all three



PI_1.c -> Using
atomic section

PI_2.c -> Using
array
(bad cache-hit
ratio)

PI_3.c -> Using
padded array



Varying the scheduling strategy

PI_Loop_1.c

```
34 int main (int argc , char ** argv)
35 {
36     int i;
37     unsigned int num_threads = atoi(argv[1]);
38     double x, pi, sum = 0.0;
39     double start_time, run_time;
40
41     step = 1.0/(double) num_steps;
42     sum = 0.0;
43     omp_set_num_threads(num_threads);
44     start_time = omp_get_wtime();
45 #pragma omp parallel
46 {
47     //#pragma omp single
48     //      printf(" num_threads = %d",omp_get_num_threads());
49
50 #pragma omp for reduction(+:sum) schedule(static,CHUNK)
51     for (i=1;i<= num_steps; i++){
52         x = (i-0.5)*step;
53         sum = sum + 4.0/(1.0+x*x);
54     }
55 }
56
57     pi = step * sum;
```

PI_Loop_2.c

```
46 #pragma omp parallel
47 {
48     double x;
49     //#pragma omp single
50     //      printf(" num_threads = %d",omp_get_num_threads());
51
52 #pragma omp for reduction(+:sum) schedule(static,CHUNK)
53     for (i=1;i<= num_steps; i++){
54         x = (i-0.5)*step;
55         sum = sum + 4.0/(1.0+x*x);
56     }
57 }
58
59     pi = step * sum;
```

In the 2nd program, each thread given its own local copy of variable 'x' - this speeds up the program by about 10 times




Different types of scheduling we tried


- Default - **STATIC**. (each thread is assigned a set of loop iterations in round-robin fashion)
- Not always optimal (for e.g. when different iterations take different amounts of time)
- In dynamic, one iteration to each thread. When the thread finishes, it will be assigned the next iteration that hasn't been executed yet.
- However, there is some overhead to dynamic scheduling. (After each iteration, the threads must stop and receive a new value of the loop variable to use for its next iteration.)
- We can split the difference between static and dynamic scheduling by using *chunks* in a dynamic schedule.

PI_Loop_2.c

Dynamic Scheduling, experimented with different chunk sizes



Comparison of Static and dynamic scheduling






Analysis of static and dynamic Scheduling

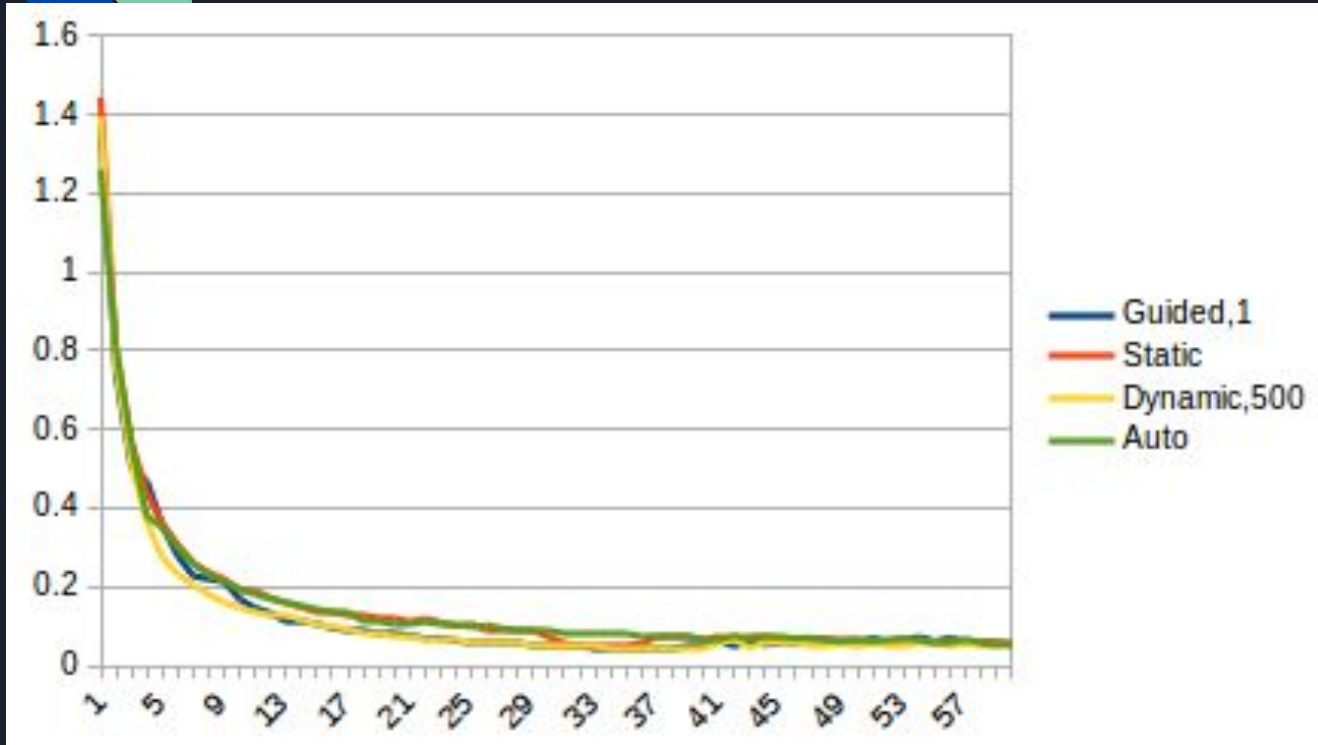
- Increasing the chunk size makes the scheduling more static.
- Decreasing the chunk size makes it more dynamic in nature.

PI_Loop_2.c -> guided scheduling



Very little change observed as we changed the chunk size for guided scheduling.


PI_Loop_2.c -> which is best?



In this program, all iterations take the same time so dynamic scheduling is not useful and it adds unnecessary overhead also.


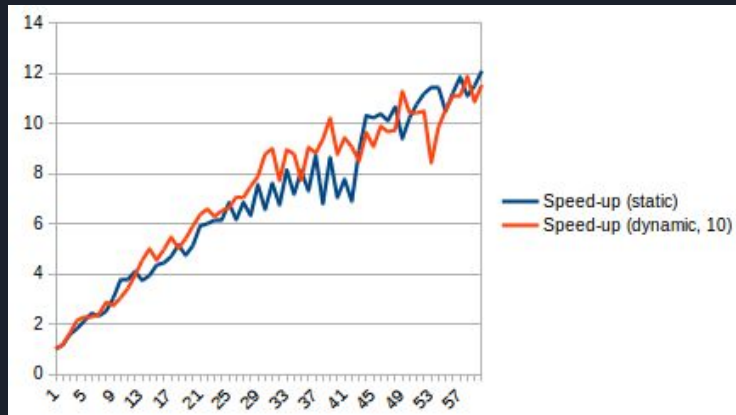

Auto/static scheduling give the optimal time.

Mandel.c

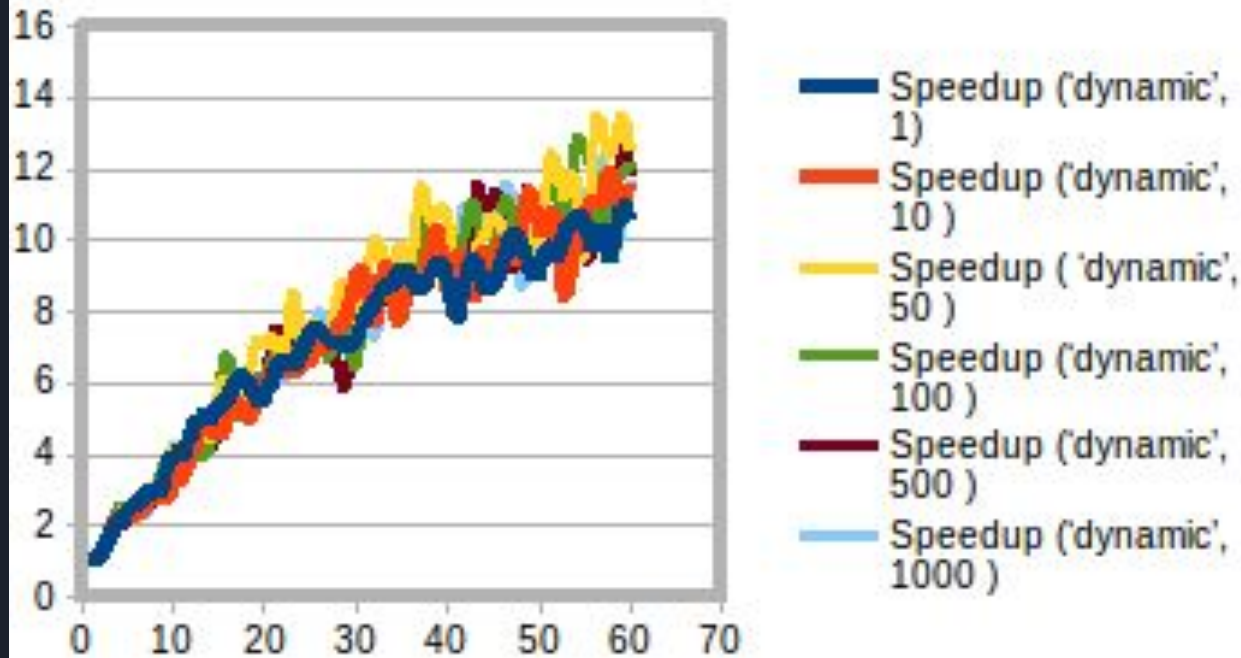


Optimal :
60 threads

Mandel.c



Mandel.c



We observed that, in this program, up till a particular chunk size (50) the speed up, in general, increased with increase in chunk size. After that the speed up remained almost same.

Mandel.c (guided scheduling)

Guided Schedule, experimented with different chunk sizes

